



# IMAGE CLASSIFICATION PROJECT

SUBMITTED BY:

NAINCY JOSHI

Problem statement

E-commerce are the biggest platform of data collection .

Collecting and sorting the data from the biggest E-commerce site is a challenge. Product images provide a better first impression.

According to a survey, more than **\*\*63 percent \*\*** of consumers say that good product images are more important than product descriptions.

For an *e-commerce platform*, good quality product images are instrumental in convincing shoppers to buy. Product images can help shoppers to get a better virtual “feel” about the product and engage on a deeper level. Collection of over 200 product images under Apparel category. Two gender types women and men under Apparel.

## Importing the dataset:

- The image classification dataset been scrapped using selenium through E-commerce website.
- The dataset has been uploaded on the google drive.
- We have mounted the google colab with the google drive so that we can connect the dataset through the drive.

```
In [25]: import warnings
warnings.filterwarnings("ignore")

import os
import glob
import shutil
```

```
In [26]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
In [1]: !nvidia-smi
```

Sun Nov 28 11:27:11 2021

NVIDIA-SMI 495.44 Driver Version: 460.32.03 CUDA Version: 11.2										
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC				
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.				
							MIG M.			
0	Tesla K80	Off	00000000:00:04:0	Off	0	0				
N/A	66C	P8	36W / 149W	0MiB / 11441MiB	0%	Default				
							N/A			

Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	
ID	ID					Usage	
No running processes found							

## Importing the Library

- We have imported the important libraries and then stored the path of the train data and the test data.
- The train data path has been stored with the attribute `train_path` and the test data been stored on the `test_path`
- We will be using Inception V3 for image classification.

```
# import the libraries
from tensorflow.keras.layers import Input, Lambda, Dense, Flatten
from tensorflow.keras.models import Model
from tensorflow.keras.applications.inception_v3 import InceptionV3
# from Keras.applications.vgg16 import VGG16
from tensorflow.keras.applications.inception_v3 import preprocess_input
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img, img_to_array
from tensorflow.keras.models import Sequential

import numpy as np
from glob import glob
import tensorflow as tf
from IPython.display import Image, display
```

```
# resize all the image to same size
Image_SIZE = [224,224]
train_path = "/content/drive/MyDrive/Hindu paper/Dataset/Train_data"
test_path = "/content/drive/MyDrive/Hindu paper/Dataset/Test_data"

# Import the Inception V3 library as shown below and ass preprocesing layer to the front of VGG
# Here we will be using imagenet weights

inception= InceptionV3(input_shape=Image_SIZE + [3], weights='imagenet', include_top=False)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_v3/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5
87916544/87910968 [=====] - 1s 0us/step
87924736/87910968 [=====] - 1s 0us/step
```

## Inception V3

1. Factorization into Smaller Convolutions
- 2. Spatial Factorization into Asymmetric Convolutions
- 3. Utility of Auxiliary Classifiers
- 4. Efficient Grid Size Reduction
- 5. The total loss used by the inception net during training.

**total\_loss = real\_loss + 0.3 \*  
aux\_loss\_1 + 0.3 \* aux\_loss\_2**

```

In [6]: # Don't train the existing weights
        for layer in inception.layers:
            layer.trainable = False

In [7]: # useful for g
        folders = glob('/content/drive/MyDrive/Hindu paper/Dataset/Train_data/*')
        folders

Out[7]: ['/content/drive/MyDrive/Hindu paper/Dataset/Train_data/Trousers',
         '/content/drive/MyDrive/Hindu paper/Dataset/Train_data/Saree',
         '/content/drive/MyDrive/Hindu paper/Dataset/Train_data/Jeans']

In [8]: # our layers- we can add more if we want
        x = Flatten()(inception.output)

In [9]: prediction = Dense(len(folders),activation='softmax')(x)

In [10]: # create a model Object
        model = Model(inputs=inception.input,outputs=prediction)

In [11]: # View the structure of the model
        model.summary()

Model: "model"

```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 224, 224, 3)]	0	[]
conv2d (Conv2D)	(None, 111, 111, 32)	864	['input_1[0][0]']
batch normalization (BatchNorm)	(None, 111, 111, 32)	96	['conv2d[0][0]']

The first layer in this network, `tf.keras.layers.Flatten`, transforms the format of the images from a two-dimensional array (of 28 by 28 pixels) to a one-dimensional array (of  $28 * 28 = 784$  pixels). Think of this layer as unstacking rows of pixels in the image and lining them up. This layer has no parameters to learn; it only reformats the data.

After the pixels are flattened, the network consists of a sequence of two `tf.keras.layers.Dense` layers. These are densely connected, or fully connected, neural layers. The

first Dense layer has 128 nodes (or neurons). The second (and last) layer returns a logits array with length of 10. Each node contains a score that indicates the current image belongs to one of the 10 classes.

## Pre-Processing

- **Softmax** extends this idea into a multi-class world. That is, Softmax assigns decimal probabilities to each class in a multi-class problem. Those decimal probabilities must add up to 1.0. This additional constraint helps training converge more quickly than it otherwise would.
- Softmax might produce the following likelihoods of an image belonging to a particular class:
- `Model . Summary()` will provide the details of the project shape and image size and many more

# Model Summary

```
In [11]: # View the structure of the model
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 224, 224, 3)	0	input_1[0][0]
conv2d (Conv2D)	(None, 111, 111, 32)	864	conv2d[0][0]
batch_normalization (Batch Normalization)	(None, 111, 111, 32)	96	batch_normalization[0][0]
activation (Activation)	(None, 111, 111, 32)	0	activation[0][0]
conv2d_1 (Conv2D)	(None, 109, 109, 32)	9216	conv2d_1[0][0]
batch_normalization_1 (Batch Normalization)	(None, 109, 109, 32)	96	batch_normalization_1[0][0]
activation_1 (Activation)	(None, 109, 109, 32)	0	activation_1[0][0]
conv2d_2 (Conv2D)	(None, 109, 109, 64)	18432	conv2d_2[0][0]
batch_normalization_2 (Batch Normalization)	(None, 109, 109, 64)	192	batch_normalization_2[0][0]
activation_2 (Activation)	(None, 109, 109, 64)	0	activation_2[0][0]
max_pooling2d (MaxPooling2D)	(None, 54, 54, 64)	0	max_pooling2d[0][0]

# Model Compile

```
In [12]: # tell the model what cost and optimization method to use
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

```
In [13]: # ...
```

For this project, we have choose the [tf.keras.optimizers.Adam](#) optimizer and [tf.keras.losses.CategoricalCrossentropy](#) loss function. To view training and validation accuracy for each training epoch, pass the metrics argument to [Model.compile](#).

[Loss function](#) —This measures how accurate the model is during training. You want to minimize this function to "steer" the model in the right direction.

[Optimizer](#) —This is how the model is updated based on the data it sees and its loss function.

[Metrics](#) —Used to monitor the training and testing steps. The following example uses *accuracy*, the fraction of the images that are correctly classified.

## Image Data Generator

```
In [13]: # use the ImageDataGenerator to import the images from the dataset
from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rescale=1./255,
                                   shear_range=0.2,
                                   zoom_range=0.2,
                                   horizontal_flip=True)
test_datagen = ImageDataGenerator(rescale = 1./255)

In [14]: # make sure that we provide the same target size as initiated for the image size
training_set=train_datagen.flow_from_directory('/content/drive/MyDrive/Hindu paper/Dataset/Train_d',
                                                target_size=(224,224),
                                                batch_size=32,
                                                class_mode='categorical')

Found 164 images belonging to 3 classes.

In [15]: training_set.class_indices

Out[15]: {'Jeans': 0, 'Saree': 1, 'Trousers': 2}
```



A data generator can also be used to specify the validation dataset and the test dataset. Often, a separate Image Data Generator instance is used that have the same pixel scaling configuration.

- As we can see that there is three classes in the data set. The classes are Jeans , saree and the Trouser.
- We have set the imaze size as 244\*244.
- We took shear imaze and the zoom imaze as 20%.

## Fitting the model

```
In [17]: #fit the model
          #Run the cell .It will take some time to execute
          r=model.fit_generator(training_set,
                                validation_data=test_set,
                                epochs=10,
                                steps_per_epoch=len(training_set),
                                validation_steps=len(test_set)
                                )

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:7: UserWarning: `Model.fit_generator`
is deprecated and will be removed in a future version. Please use `Model.fit`, which supports gener
ators.
  import sys
Epoch 1/10
6/6 [=====] - 54s 6s/step - loss: 3.7272 - accuracy: 0.4878 - val_loss: 1.
2872 - val_accuracy: 0.8158
Epoch 2/10
6/6 [=====] - 3s 570ms/step - loss: 1.9289 - accuracy: 0.7683 - val_loss:
3.0203 - val_accuracy: 0.7368
Epoch 3/10
6/6 [=====] - 3s 463ms/step - loss: 1.3629 - accuracy: 0.8293 - val_loss:
1.5753 - val_accuracy: 0.8158
Epoch 4/10
6/6 [=====] - 3s 483ms/step - loss: 0.7134 - accuracy: 0.8963 - val_loss:
0.6317 - val_accuracy: 0.8694
```

```

6/6 [=====] - 3s 570ms/step - loss: 1.9289 - accuracy: 0.7683 - val_loss:
3.0203 - val_accuracy: 0.7368
Epoch 3/10
6/6 [=====] - 3s 463ms/step - loss: 1.3629 - accuracy: 0.8293 - val_loss:
1.5753 - val_accuracy: 0.8158
Epoch 4/10
6/6 [=====] - 3s 483ms/step - loss: 0.7134 - accuracy: 0.8963 - val_loss:
0.6317 - val_accuracy: 0.8684
Epoch 5/10
6/6 [=====] - 3s 480ms/step - loss: 0.5480 - accuracy: 0.8720 - val_loss:
0.5812 - val_accuracy: 0.8684
Epoch 6/10
6/6 [=====] - 3s 467ms/step - loss: 0.3523 - accuracy: 0.9268 - val_loss:
0.8165 - val_accuracy: 0.8684
Epoch 7/10
6/6 [=====] - 3s 475ms/step - loss: 0.6169 - accuracy: 0.9146 - val_loss:
1.4527 - val_accuracy: 0.7895
Epoch 8/10
6/6 [=====] - 3s 460ms/step - loss: 0.5719 - accuracy: 0.9024 - val_loss:
0.8640 - val_accuracy: 0.8684
Epoch 9/10
6/6 [=====] - 3s 487ms/step - loss: 0.5212 - accuracy: 0.9146 - val_loss:
0.3705 - val_accuracy: 0.9211
Epoch 10/10
6/6 [=====] - 3s 490ms/step - loss: 0.2382 - accuracy: 0.9451 - val_loss:
1.4001 - val_accuracy: 0.8684

```


## Train and Test Loss and Accuracy

```

import matplotlib.pyplot as plt
plt.plot(r.history['loss'],label='train_loss')
plt.plot(r.history['val_loss'],label='val_loss')
plt.legend()
plt.show()
plt.savefig('LossVal_loss')

#Plot the accuracy
plt.plot(r.history['accuracy'],label='train_accuracy')
plt.plot(r.history['val_accuracy'],label='test_accuracy')
plt.legend()
plt.show()
plt.savefig('Accuracy over Epochs')

```



Predicting the test dataset:

```

In [22]: import numpy as np
        y_pred = np.argmax(y_pred, axis=1)

In [23]: y_pred

Out[23]: array([2, 0, 1, 2, 2, 1, 2, 0, 1, 2, 2, 1, 1, 2, 0, 0, 1, 2, 1, 2, 2, 1,
        1, 2, 1, 0, 1, 2, 0, 1, 2, 2, 2, 1, 2, 1, 0, 2])


In [24]: from tensorflow.keras.models import load_model
        from tensorflow.keras.preprocessing import image

In [67]: test_img, label = test_set.next()

In [72]: def plotImages(y_pred, label):
        """
        input:image array
        output:plot images
        """
        for idx, img in enumerate(y_pred):
            if idx <= 10:
                plt.figure(figsize=(5,5))
                plt.imshow(img)
                plt.title(img.shape)
                plt.axis=False
                plt.show()

In [74]: plotImages(test_img, label)

```



- As we can see that in the prediction , we got 2 which means its belong to the trouser as we can see in the image that it predicted right.

## Saving the model

```

In [19]: # Save the Model

        model.save('model_vgg16_ImageClassificationcomplete')
        model.save('model_vgg16_tf', save_format='tf')
        model.save('model_vgg16.h5')

INFO:tensorflow:Assets written to: model_vgg16_ImageClassificationcomplete/assets
INFO:tensorflow:Assets written to: model_vgg16_tf/assets

```

